

Global Infrastructure Section

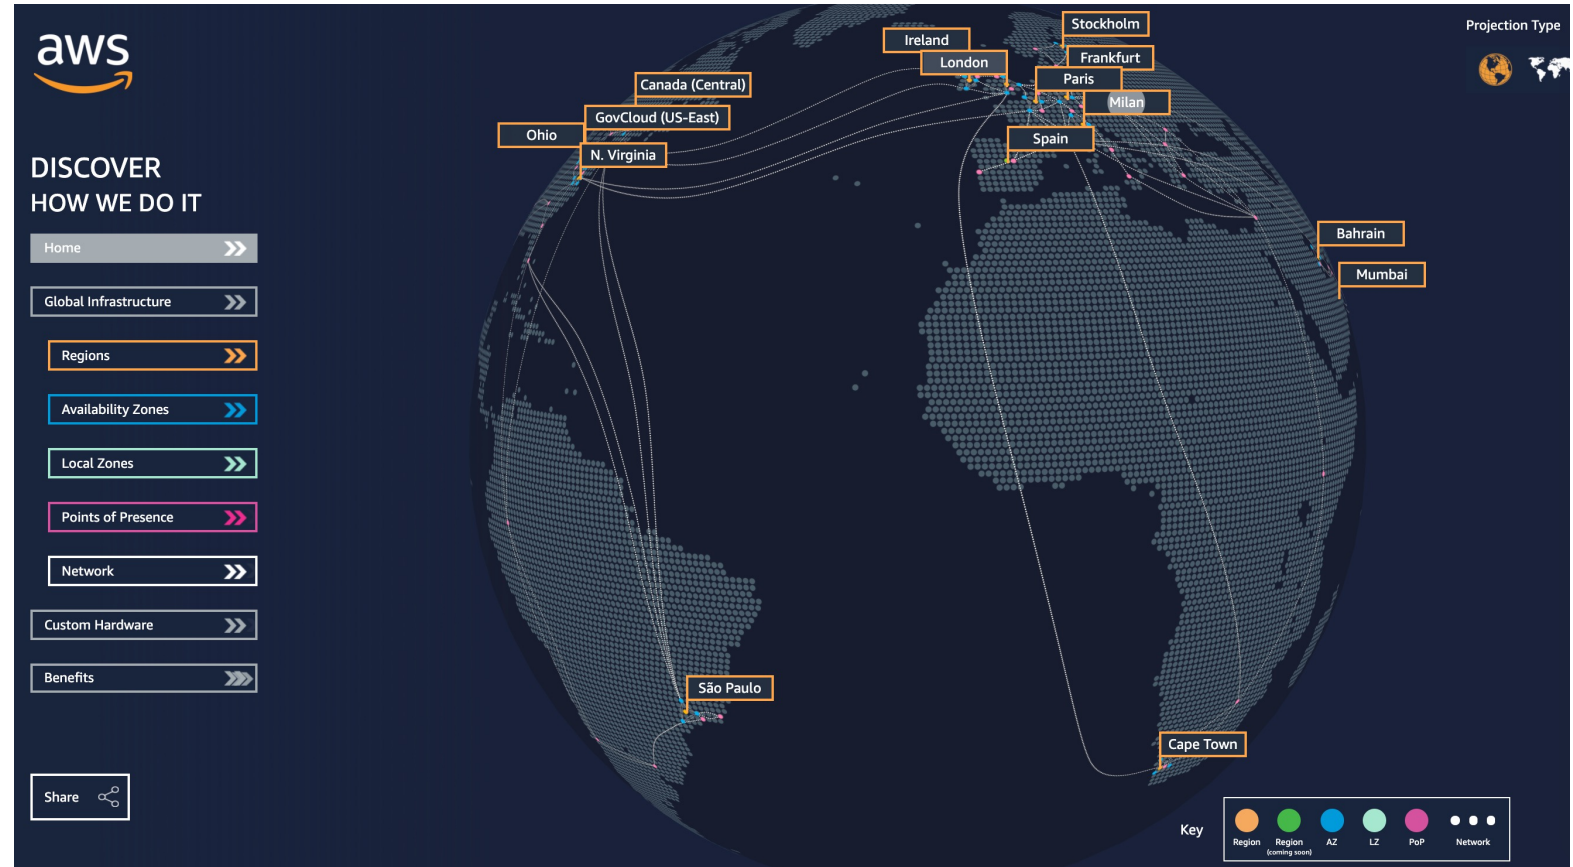


Why make a global application?

- A **global application** is an application deployed in **multiple geographies**
- On AWS: this could be **Regions** and / or **Edge Locations**
- **Decreased Latency**
 - Latency is the time it takes for a network packet to reach a server
 - It takes time for a packet from Asia to reach the US
 - Deploy your applications closer to your users to decrease latency, better experience
- **Disaster Recovery (DR)**
 - If an AWS region goes down (earthquake, storms, power shutdown, politics)...
 - You can fail-over to another region and have your application still working
 - A DR plan is important to increase the availability of your application
- **Attack protection:** distributed global infrastructure is harder to attack

Global AWS Infrastructure

- **Regions:** For deploying applications and infrastructure
- **Availability Zones:** Made of multiple data centers
- **Edge Locations (Points of Presence):** for content delivery as close as possible to users
- More at: <https://infrastructure.aws/>



Global Applications in AWS



- **Global DNS: Route 53**

- Great to route users to the closest deployment with least latency
- Great for disaster recovery strategies



- **Global Content Delivery Network (CDN): CloudFront**

- Replicate part of your application to AWS Edge Locations – decrease latency
- Cache common requests – improved user experience and decreased latency



- **S3 Transfer Acceleration**

- Accelerate global uploads & downloads into Amazon S3



- **AWS Global Accelerator:**

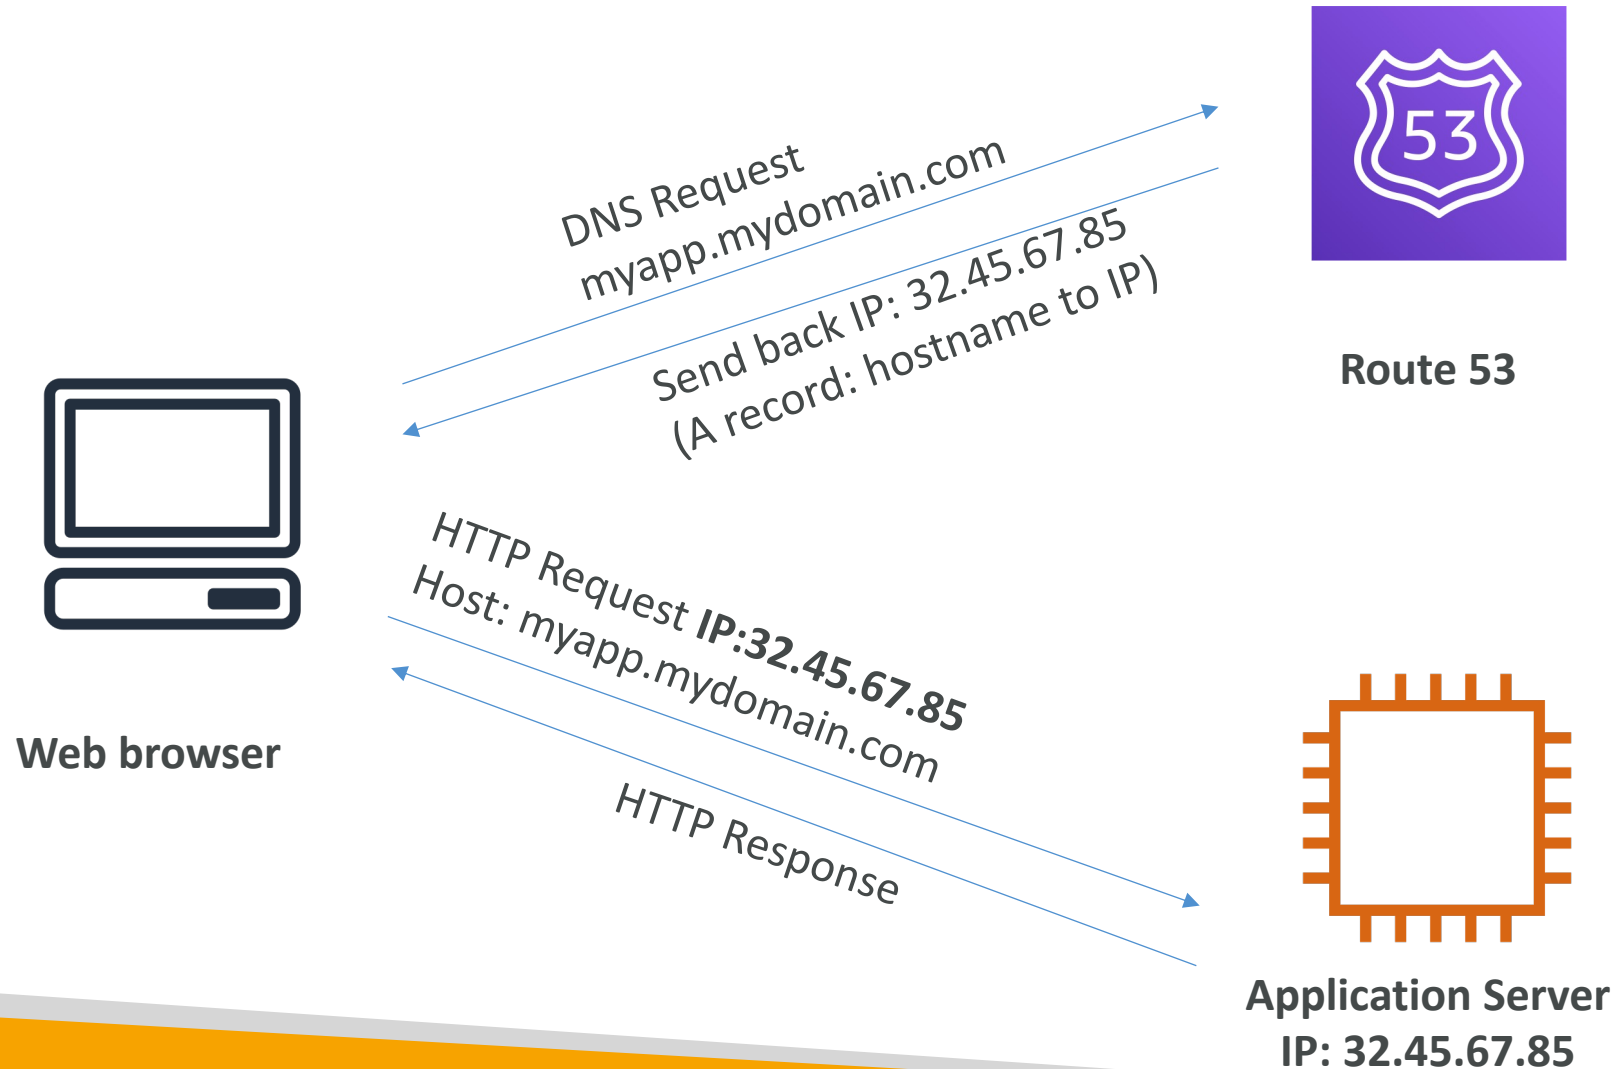
- Improve global application availability and performance using the AWS global network



Amazon Route 53 Overview

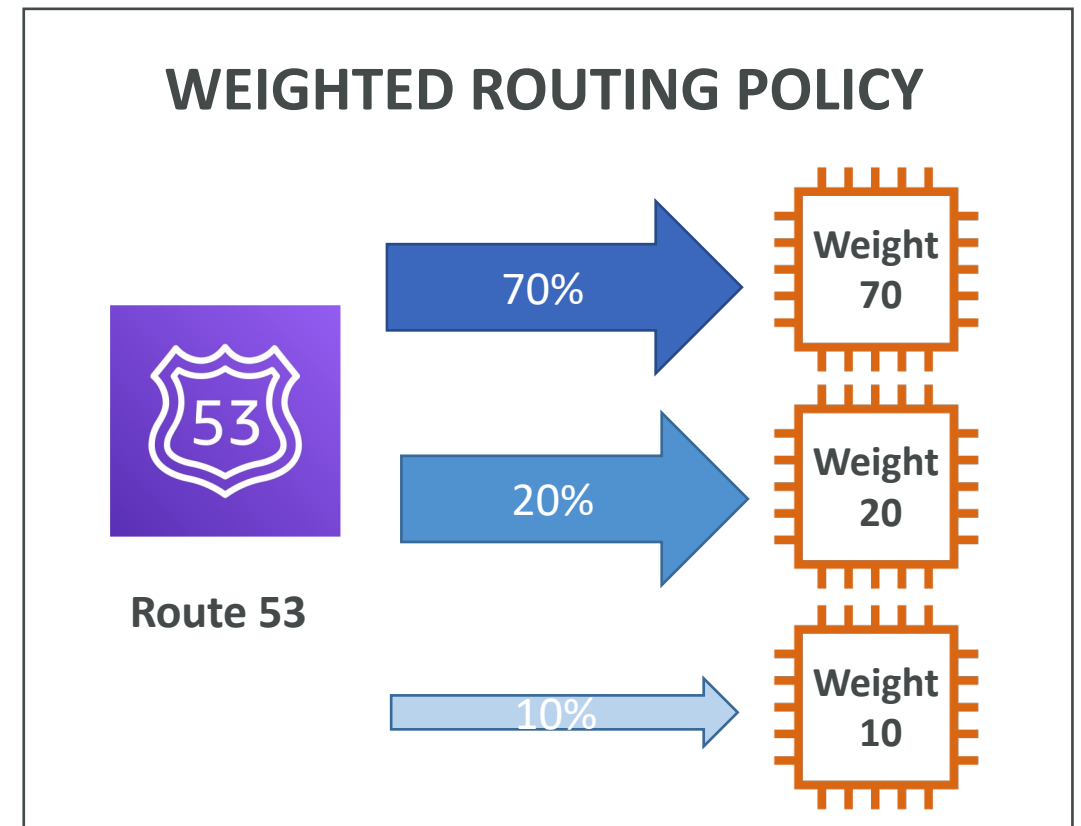
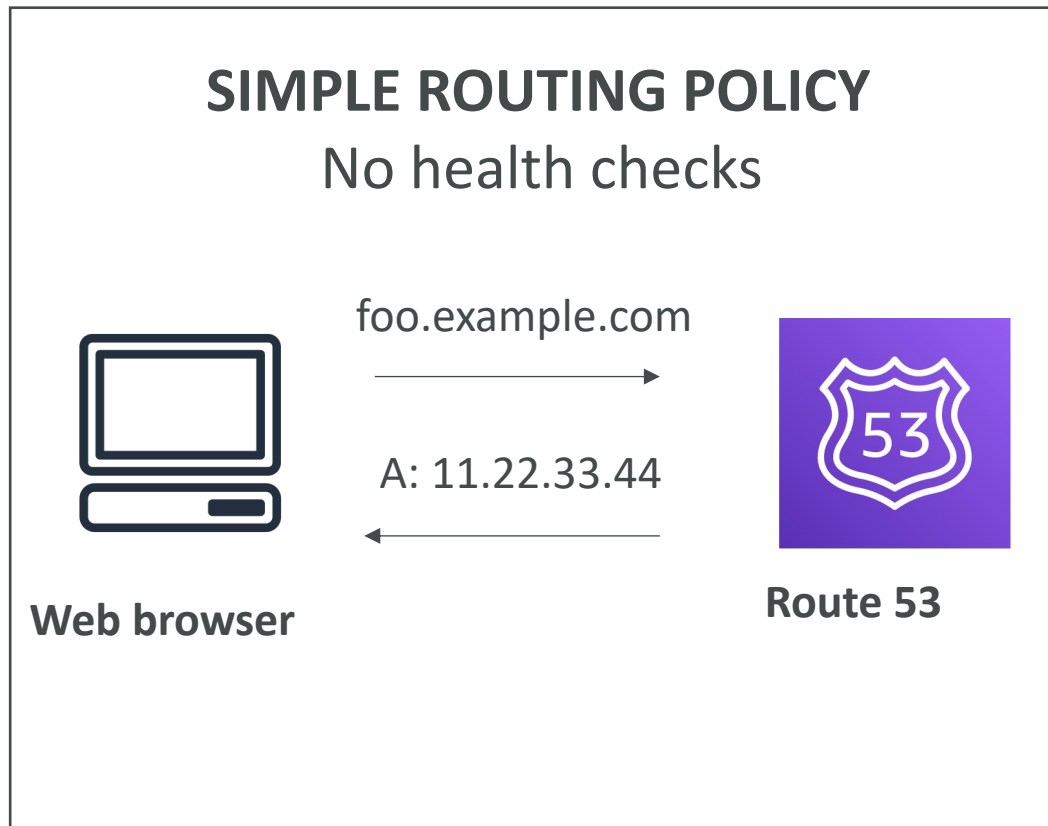
- Route53 is a Managed DNS (Domain Name System)
- DNS is a collection of rules and records which helps clients understand how to reach a server through URLs.
- In AWS, the most common records are:
 - `www.google.com => 12.34.56.78 == A record (IPv4)`
 - `www.google.com => 2001:0db8:85a3:0000:0000:8a2e:0370:7334 == AAAA IPv6`
 - `search.google.com => www.google.com == CNAME: hostname to hostname`
 - `example.com => AWS resource == Alias (ex: ELB, CloudFront, S3, RDS, etc...)`

Route 53 – Diagram for A Record



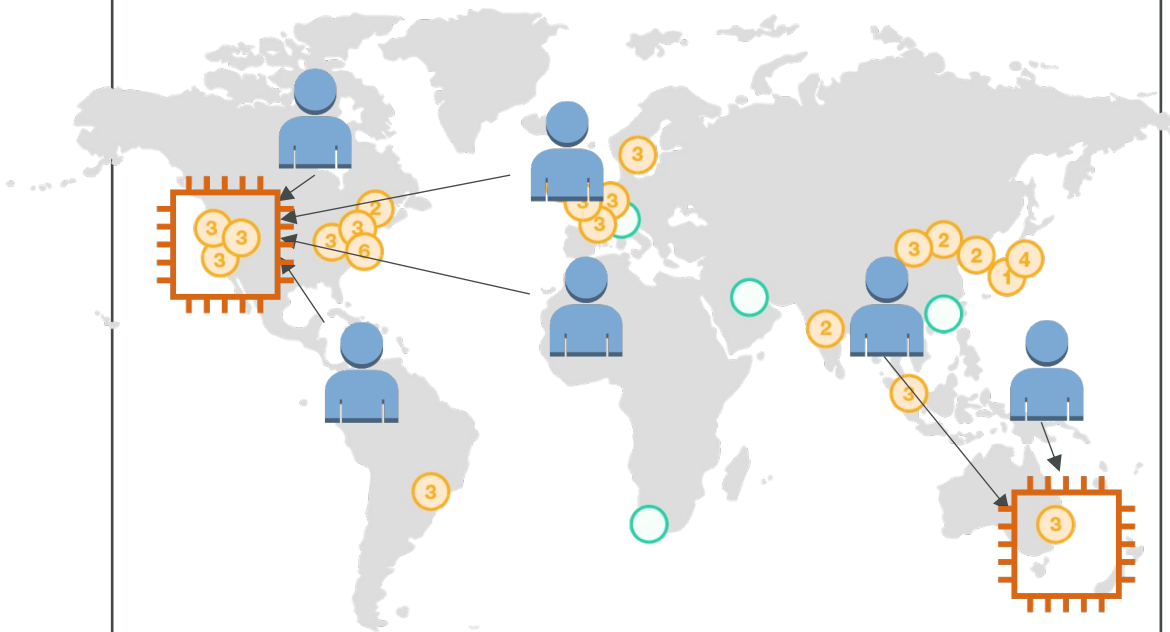
Route 53 Routing Policies

- Need to know them at a high-level for the Cloud Practitioner Exam



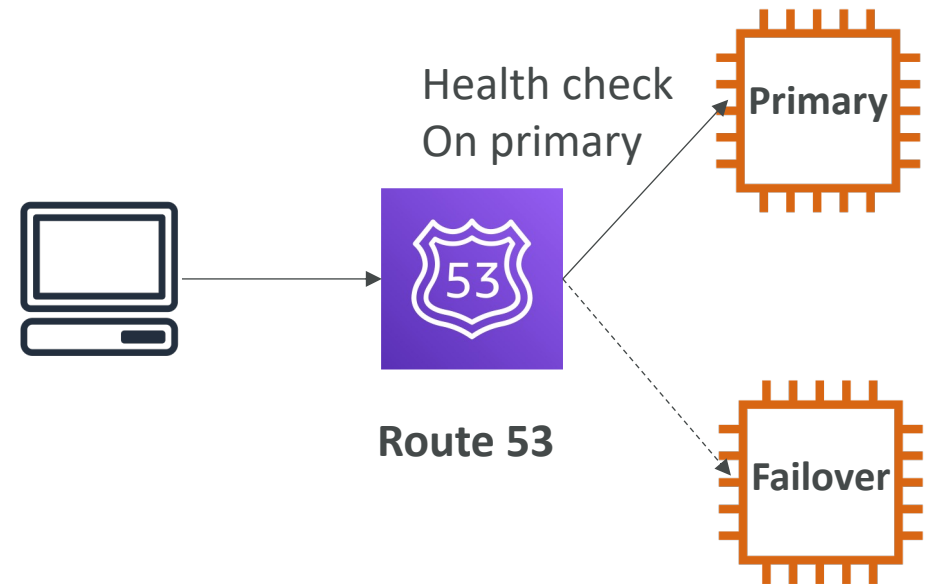
Route 53 Routing Policies

LATENCY ROUTING POLICY



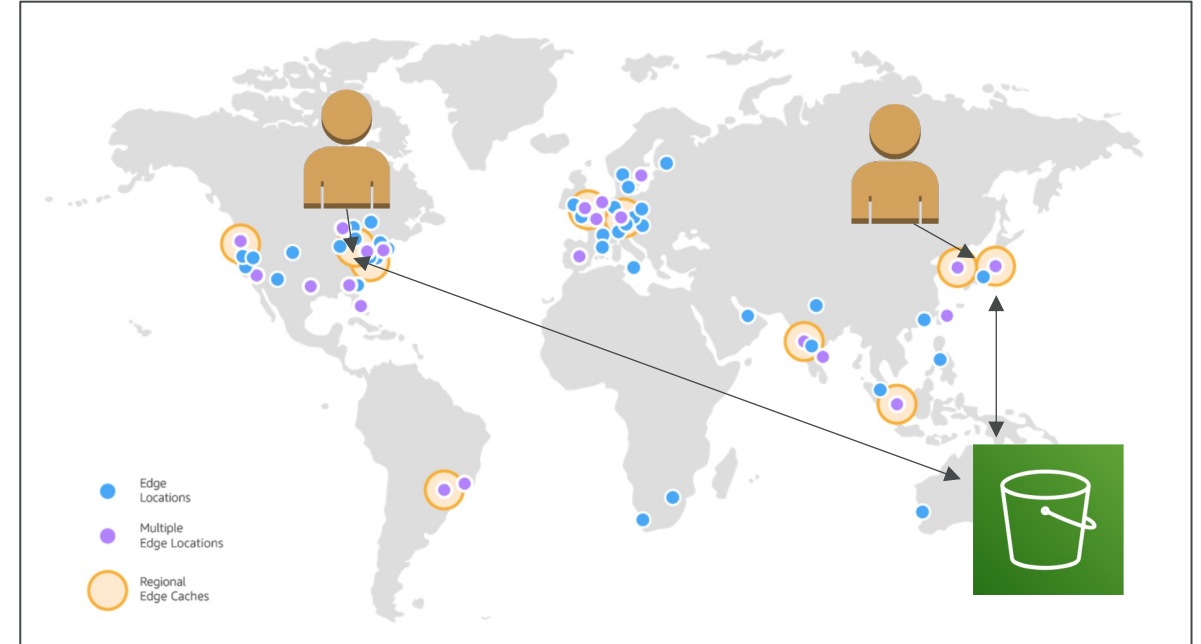
FAILOVER ROUTING POLICY

Disaster Recovery



AWS CloudFront

- Content Delivery Network (CDN)
- Improves read performance, content is cached at the edge
- Improves users experience
- 216 Point of Presence globally (edge locations)
- DDoS protection (because worldwide), integration with Shield, AWS Web Application Firewall

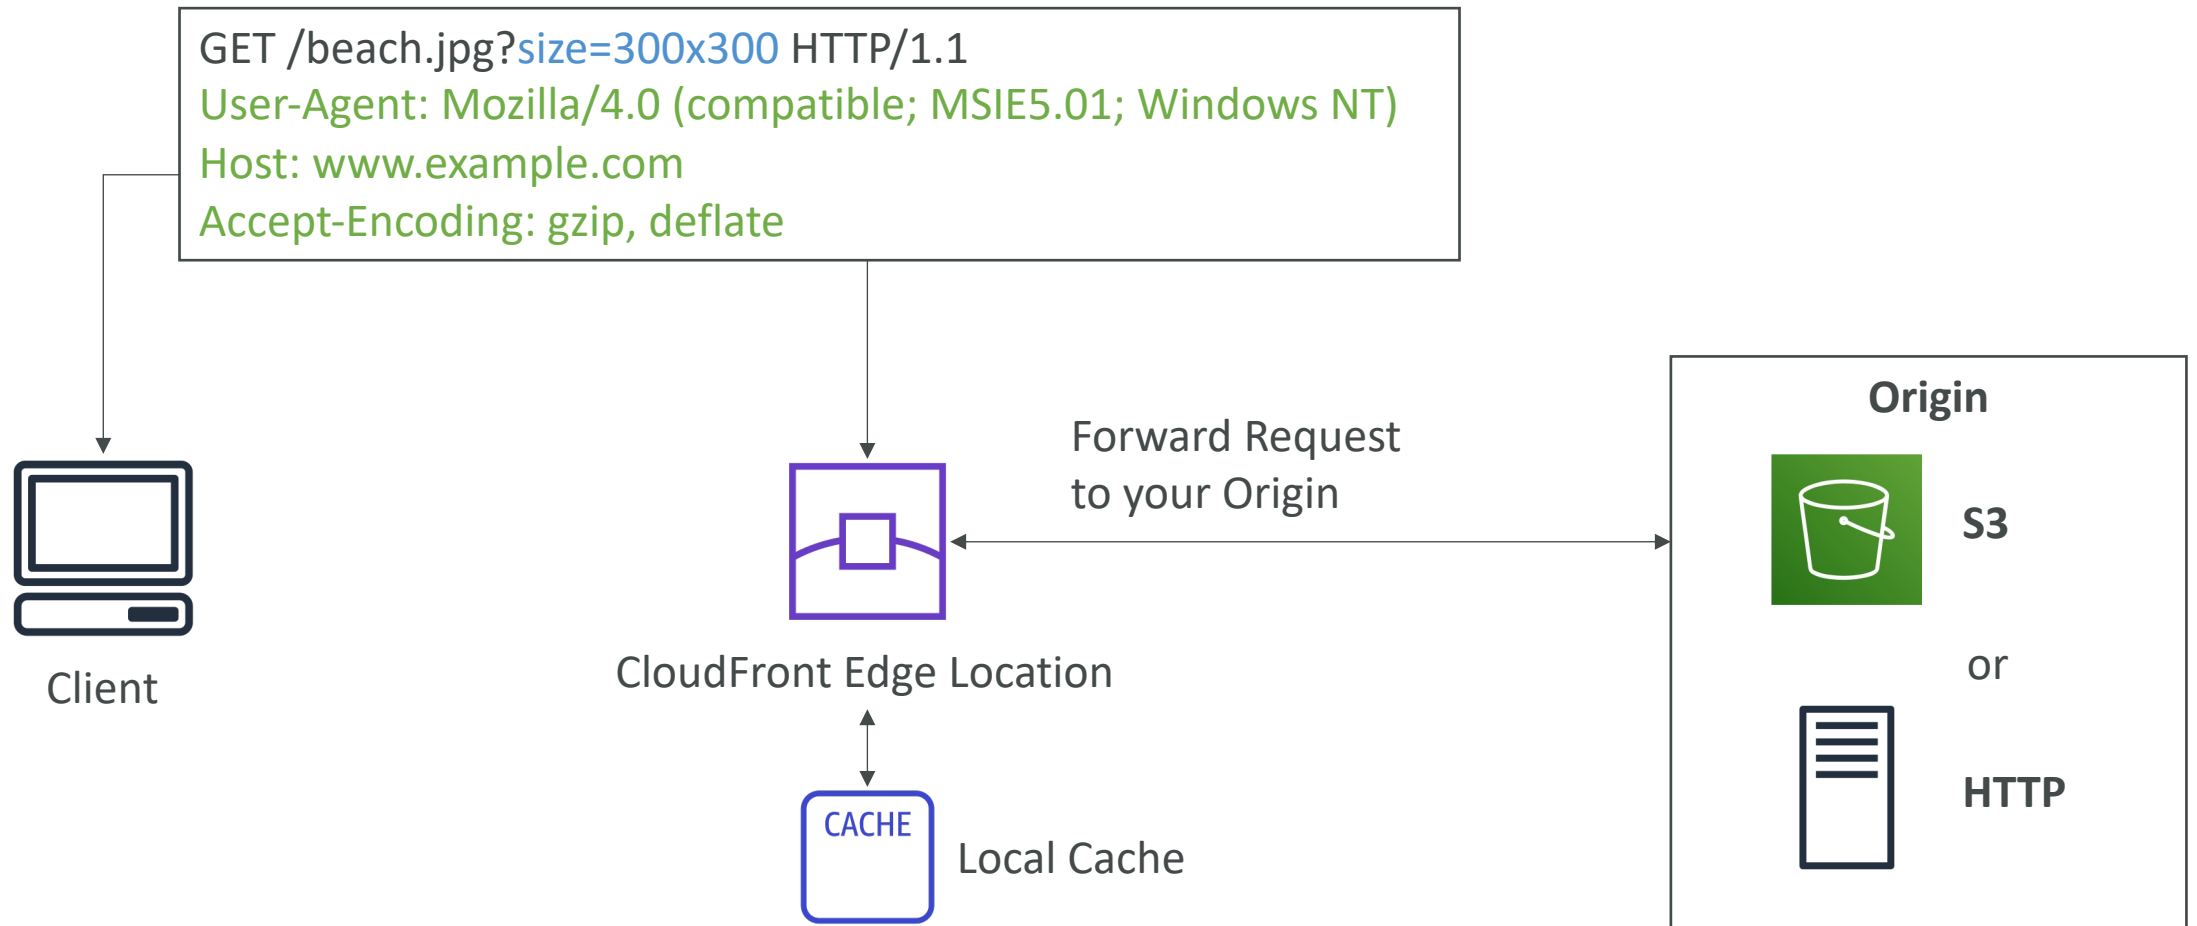


Source: <https://aws.amazon.com/cloudfront/features/?nc=sn&loc=2>

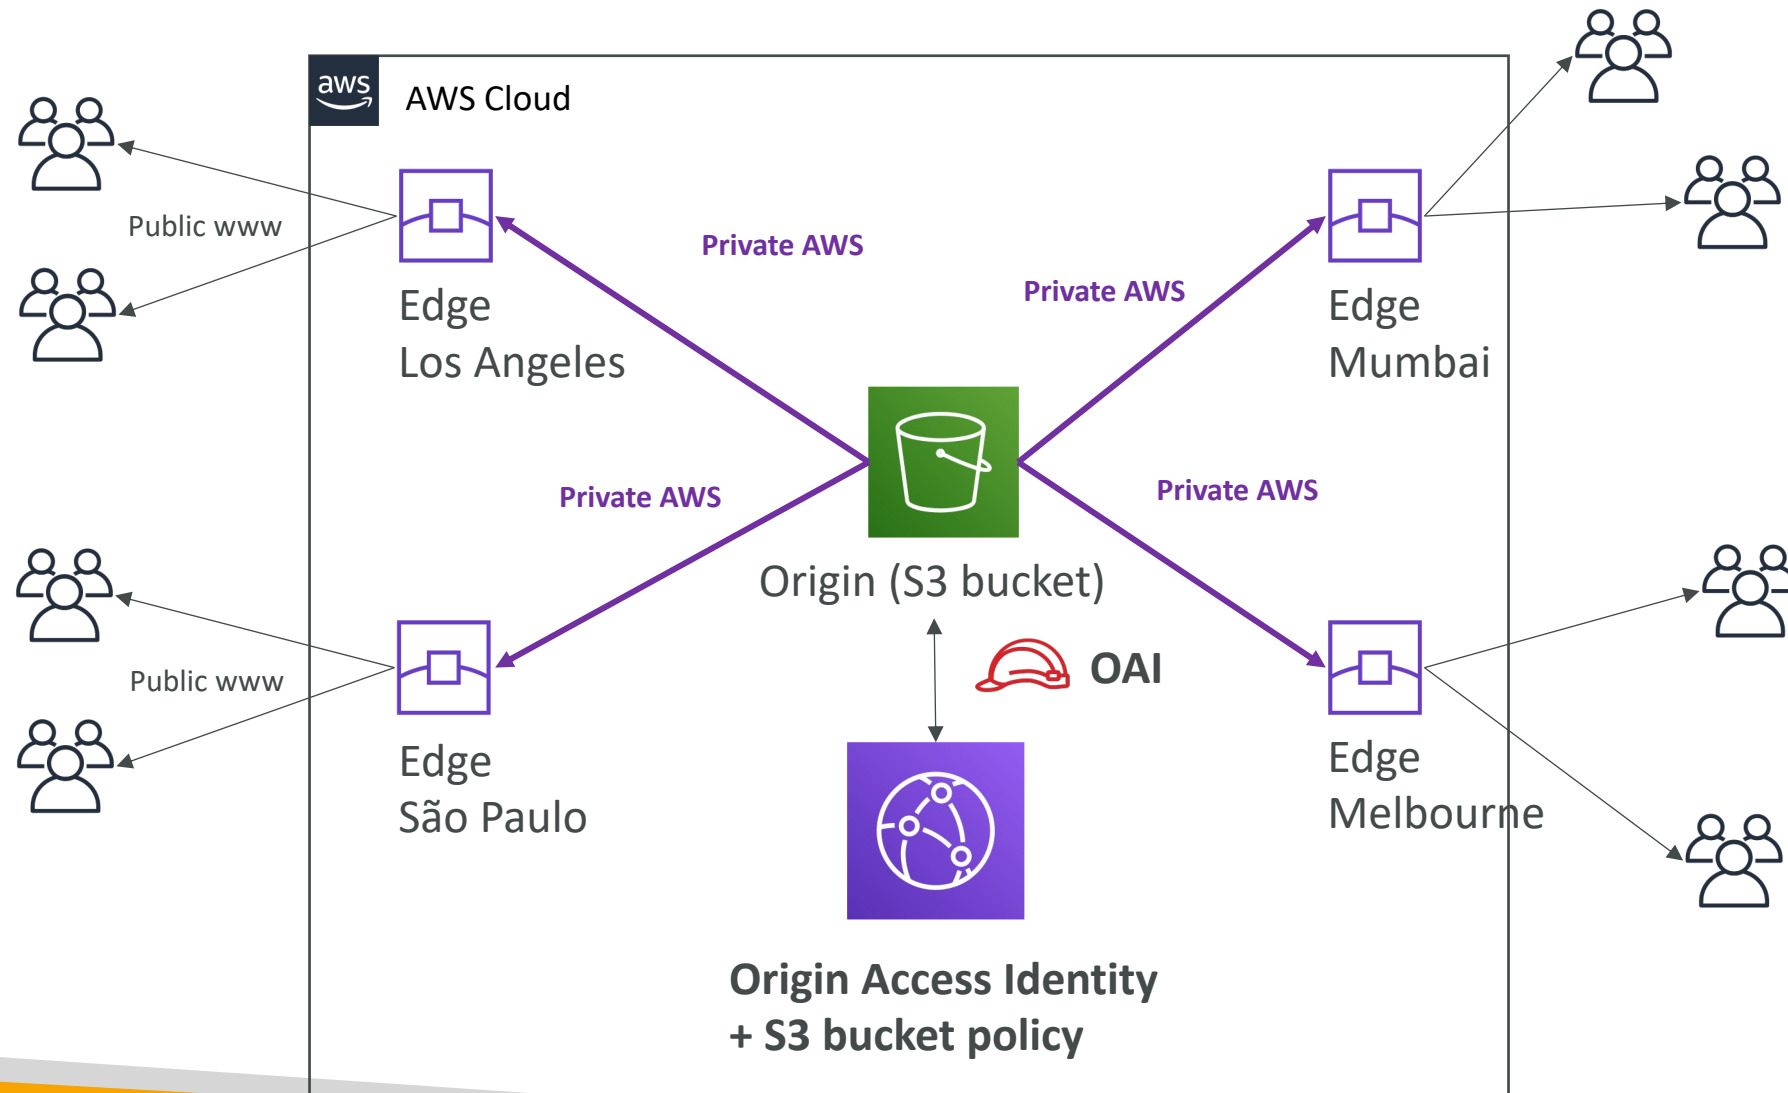
CloudFront – Origins

- S3 bucket
 - For distributing files and caching them at the edge
 - Enhanced security with CloudFront **Origin Access Identity (OAI)**
 - CloudFront can be used as an ingress (to upload files to S3)
- Custom Origin (HTTP)
 - Application Load Balancer
 - EC2 instance
 - S3 website (must first enable the bucket as a static S3 website)
 - Any HTTP backend you want

CloudFront at a high level



CloudFront – S3 as an Origin




CloudFront vs S3 Cross Region Replication

- CloudFront:
 - Global Edge network
 - Files are cached for a TTL (maybe a day)
 - Great for static content that must be available everywhere
- S3 Cross Region Replication:
 - Must be setup for each region you want replication to happen
 - Files are updated in near real-time
 - Read only
 - Great for dynamic content that needs to be available at low-latency in few regions

S3 Transfer Acceleration


- Increase transfer speed by transferring file to an AWS edge location which will forward the data to the S3 bucket in the target region



Test the tool at: <https://s3-accelerate-speedtest.s3-accelerate.amazonaws.com/en/accelerate-speed-comparision.html>

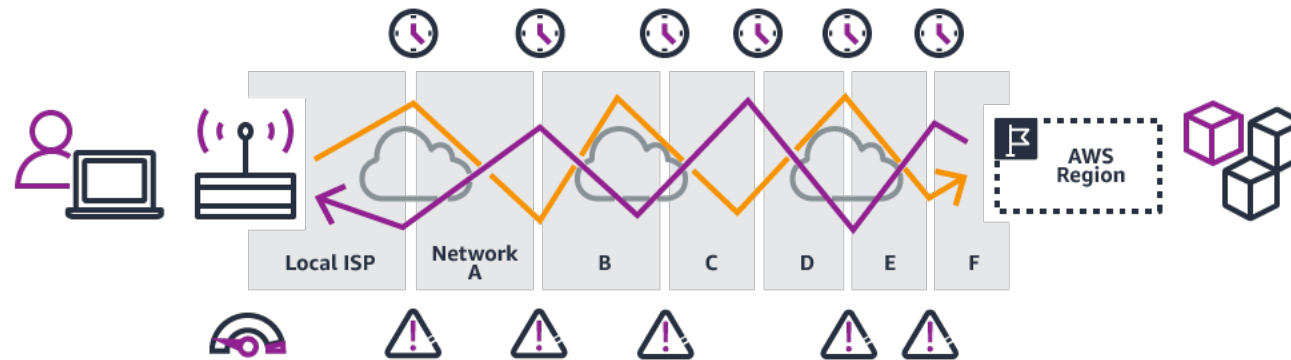
AWS Global Accelerator

- Improve global application availability and performance using the AWS global network
- Leverage the AWS internal network to optimize the route to your application (60% improvement)
- 2 **Anycast IP** are created for your application and traffic is sent through **Edge Locations**
- The Edge locations send the traffic to your application




AWS Global Accelerator

Without Global Accelerator



With Global Accelerator

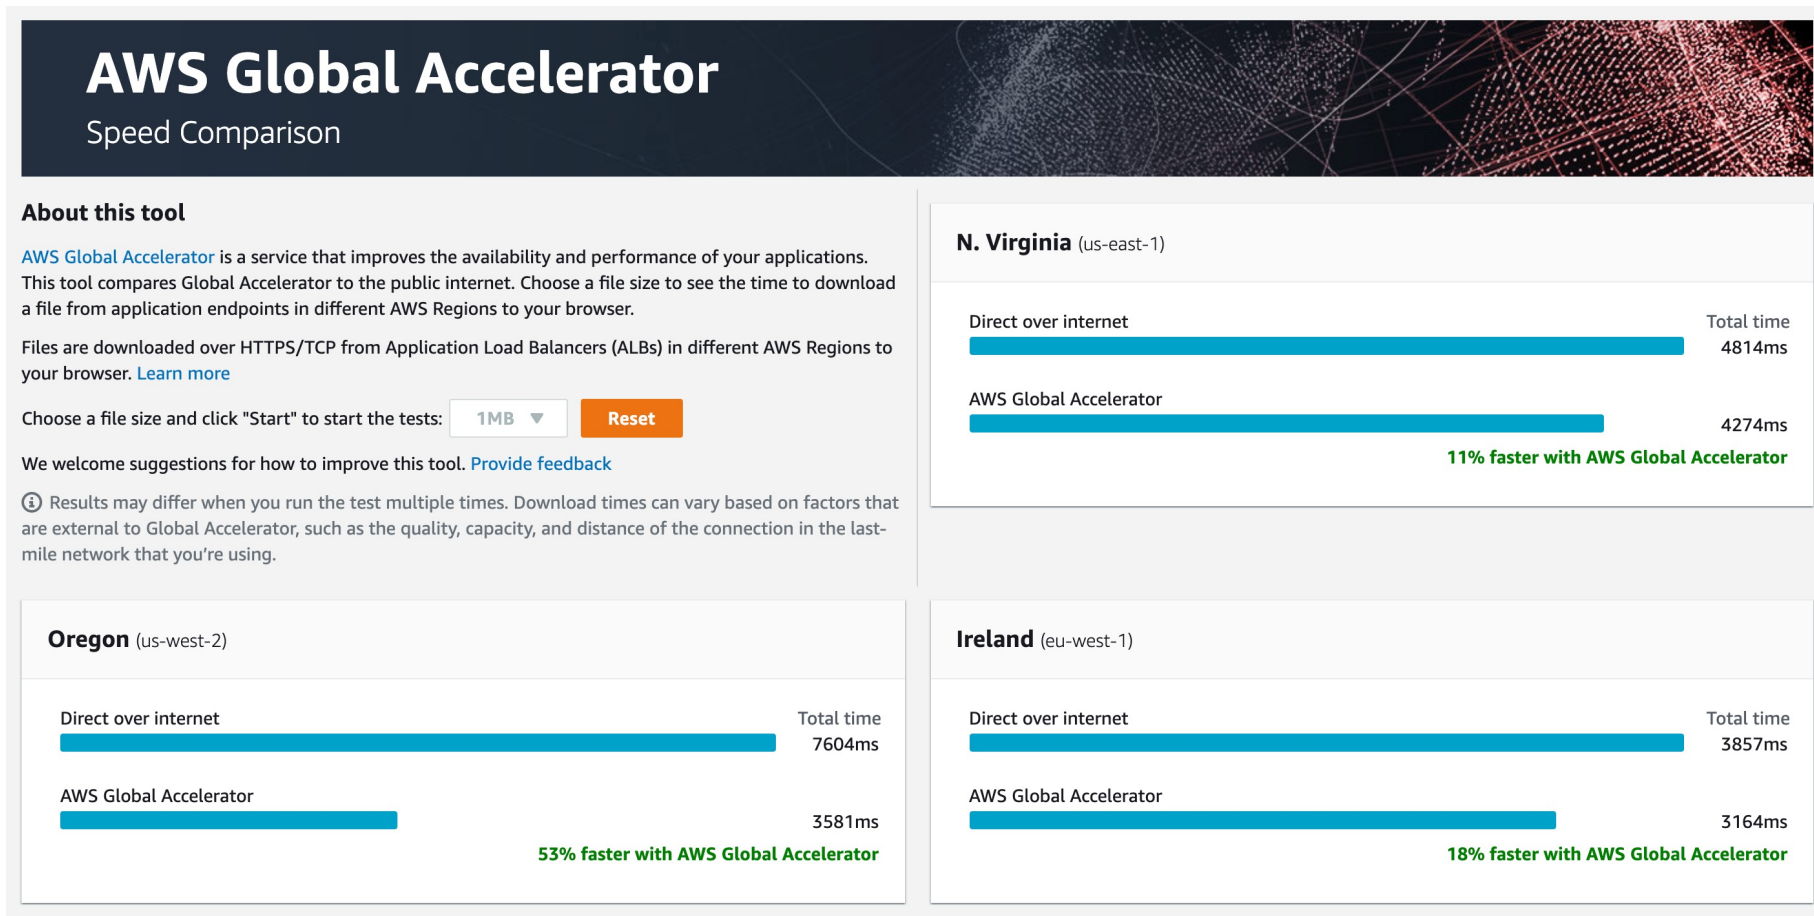


<https://aws.amazon.com/global-accelerator>


AWS Global Accelerator vs CloudFront

- They both use the AWS global network and its edge locations around the world
- Both services integrate with AWS Shield for DDoS protection.
- **CloudFront – Content Delivery Network**
 - Improves performance for your cacheable content (such as images and videos)
 - Content is served at the edge
- **Global Accelerator**
 - No caching, proxying packets at the edge to applications running in one or more AWS Regions.
 - Improves performance for a wide range of applications over TCP or UDP
 - Good for HTTP use cases that require static IP addresses
 - Good for HTTP use cases that required deterministic, fast regional failover

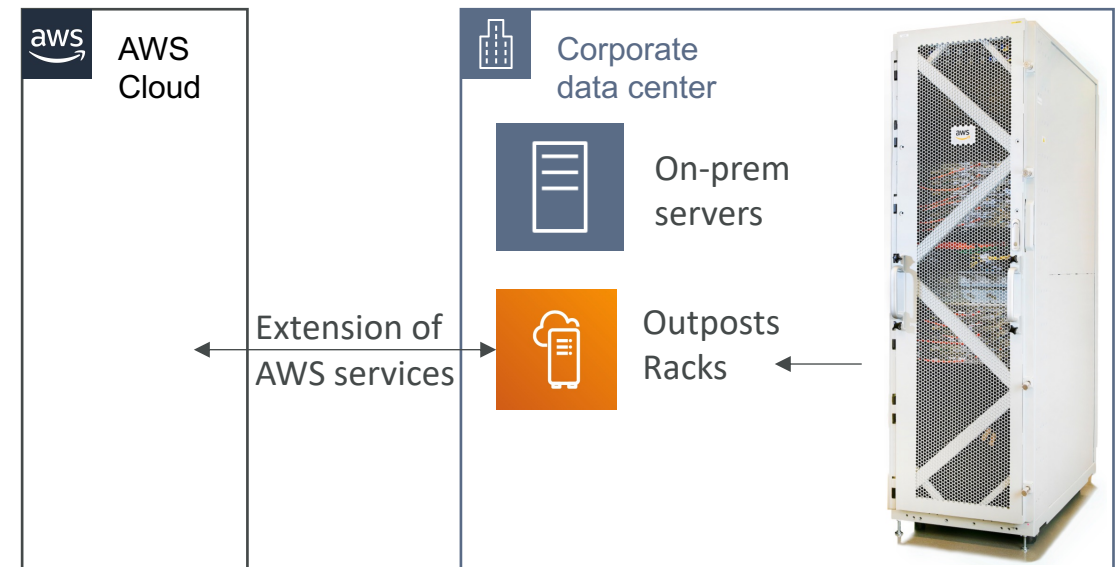
<https://speedtest.globalaccelerator.aws/#/>




AWS Outposts




- **Hybrid Cloud:** businesses that keep an on-premises infrastructure alongside a cloud infrastructure
- Therefore, two ways of dealing with IT systems:
 - One for the AWS cloud (using the AWS console, CLI, and AWS APIs)
 - One for their on-premises infrastructure
- **AWS Outposts are “server racks”** that offers the same AWS infrastructure, services, APIs & tools to build your own applications on-premises just as in the cloud
- **AWS will setup and manage “Outposts Racks”** within your on-premises infrastructure and you can start leveraging AWS services on-premises
- You are responsible for the Outposts Rack physical security






AWS Outposts

- Benefits:
 - Low-latency access to on-premises systems
 - Local data processing
 - Data residency
 - Easier migration from on-premises to the cloud
 - Fully managed service
- Some services that work on Outposts:



Amazon EC2



Amazon EBS



Amazon S3




Amazon EKS



Amazon ECS



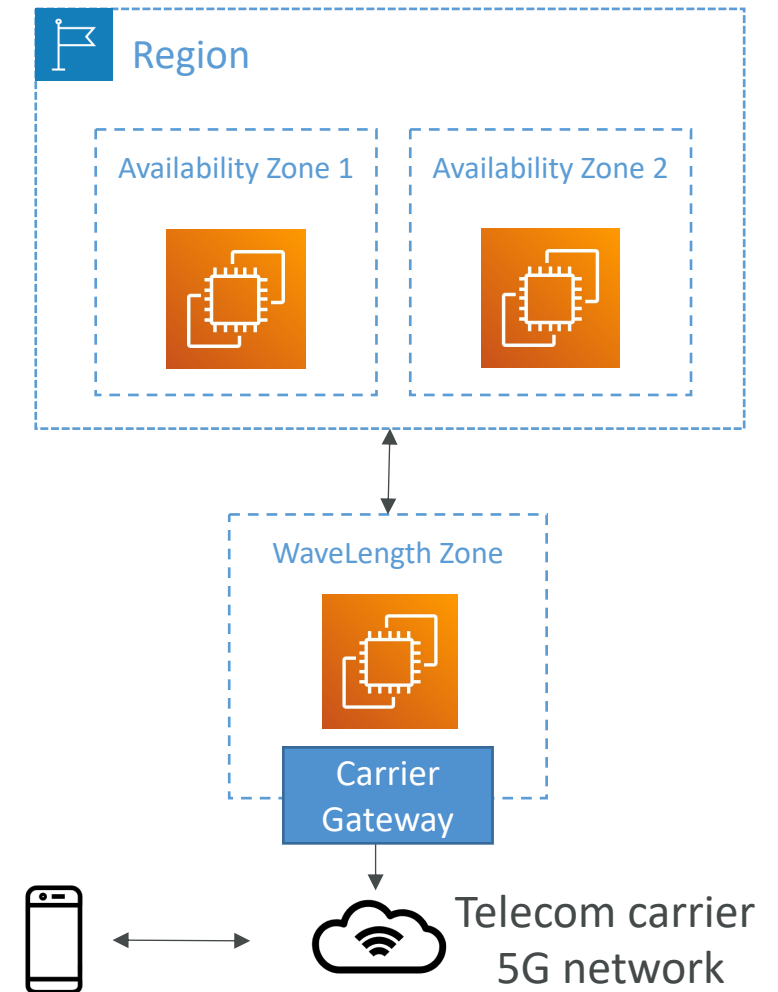
Amazon RDS



Amazon EMR

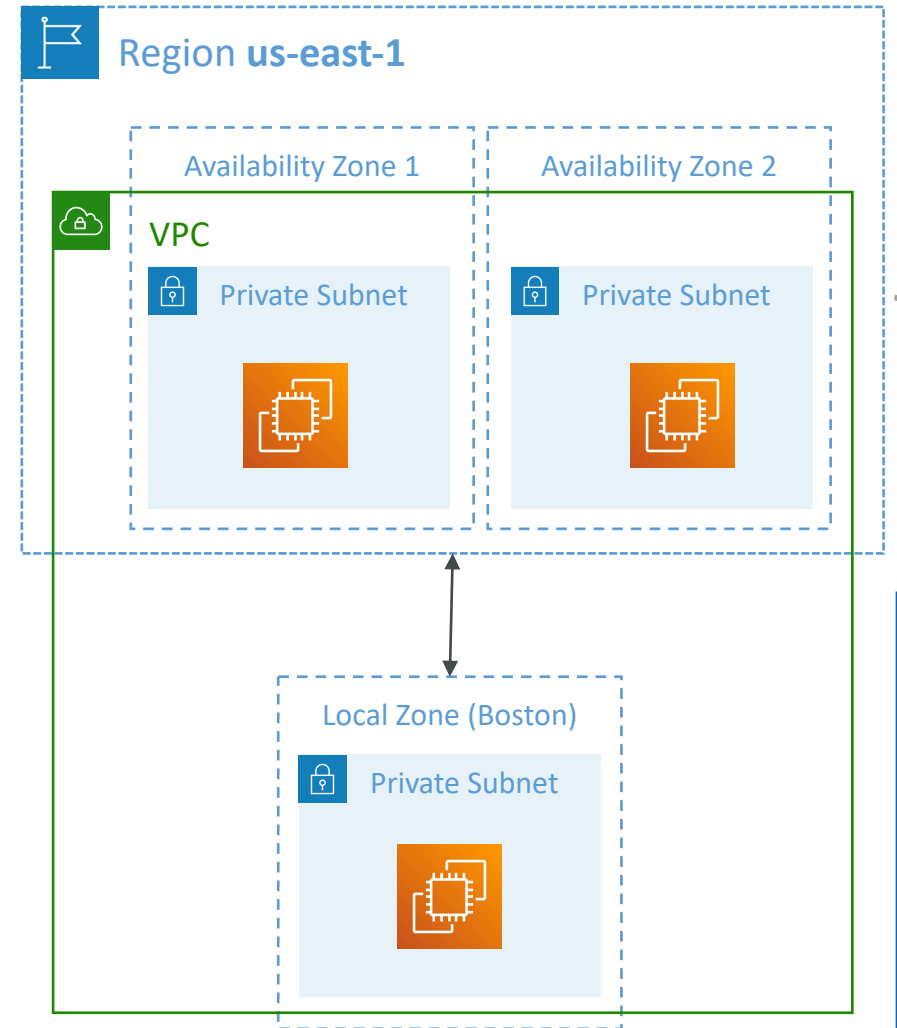
AWS WaveLength

- **WaveLength Zones** are infrastructure deployments embedded within the telecommunications providers' datacenters at the edge of the 5G networks
- Brings AWS services to the edge of the 5G networks
- Example: EC2, EBS, VPC...
- Ultra-low latency applications through 5G networks
- Traffic doesn't leave the Communication Service Provider's (CSP) network
- High-bandwidth and secure connection to the parent AWS Region
- No additional charges or service agreements
- Use cases: Smart Cities, ML-assisted diagnostics, Connected Vehicles, Interactive Live Video Streams, AR/VR, Real-time Gaming, ...



AWS Local Zones

- Places AWS compute, storage, database, and other selected AWS services **closer to end users** to run **latency-sensitive applications**
- Extend your VPC to more locations – “**Extension of an AWS Region**”
- Compatible with EC2, RDS, ECS, EBS, ElastiCache, Direct Connect ...
- Example:
 - **AWS Region:** N. Virginia (us-east-1)
 - **AWS Local Zones:** Boston, Chicago, Dallas, Houston, Miami, ...



Global Applications Architecture


Single Region, Single AZ

- ✗ High Availability
- ✗ Global Latency
- 📈 Difficulty



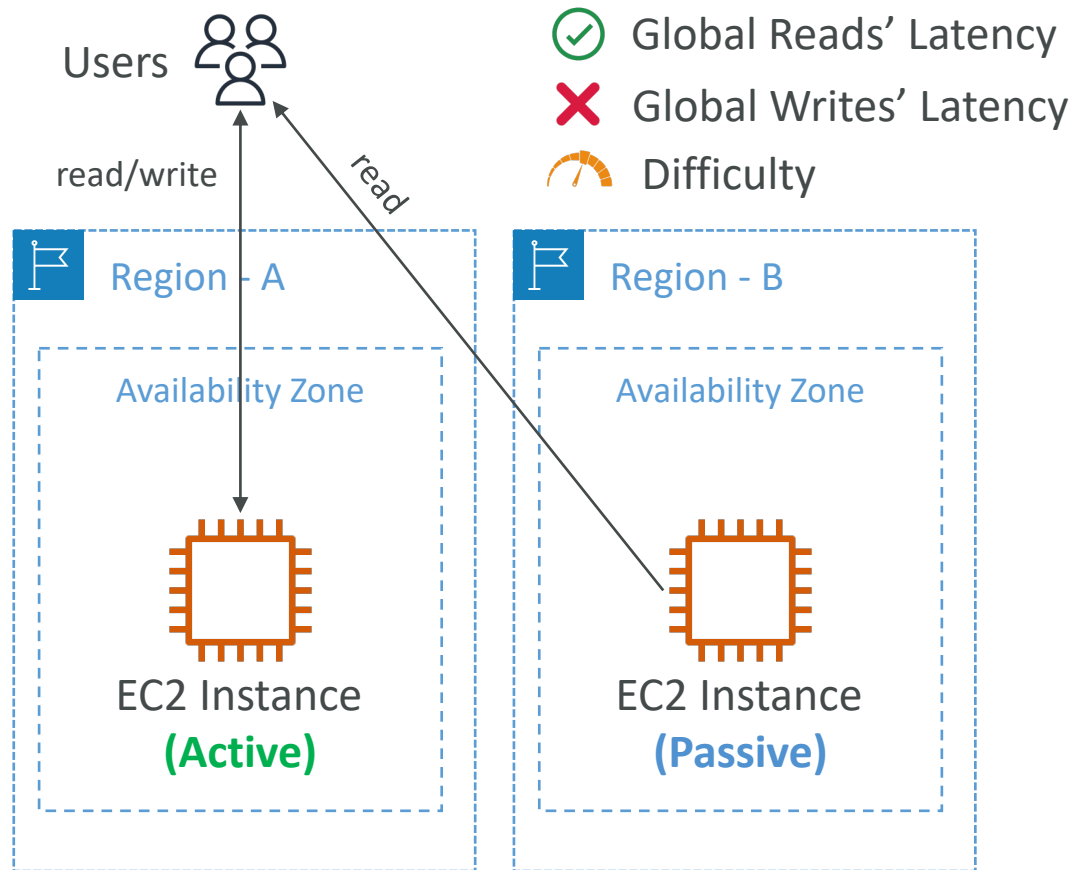
Single Region, Multi AZ

- ✓ High Availability
- ✗ Global Latency
- 📈 Difficulty

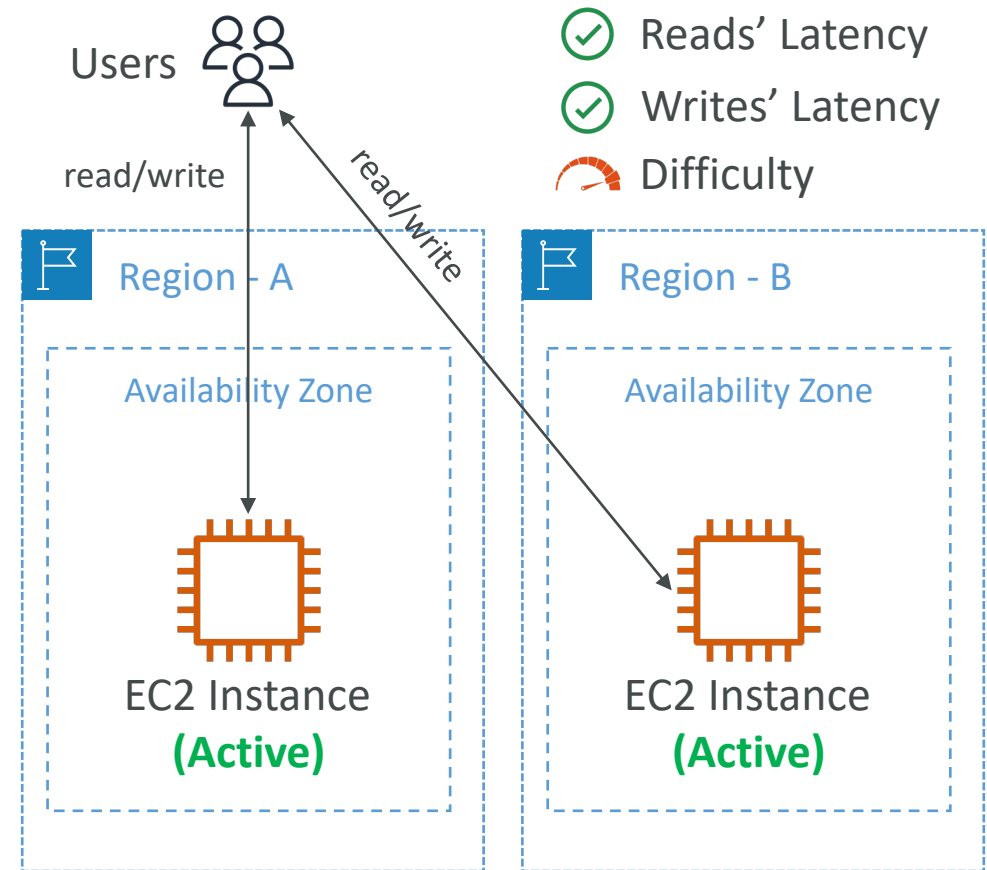


Global Applications Architecture

Multi Region, Active-Passive



Multi Region, Active-Active



Global Applications in AWS - Summary




- **Global DNS: Route 53**

- Great to route users to the closest deployment with least latency
- Great for disaster recovery strategies



- **Global Content Delivery Network (CDN): CloudFront**

- Replicate part of your application to AWS Edge Locations – decrease latency
- Cache common requests – improved user experience and decreased latency



- **S3 Transfer Acceleration**

- Accelerate global uploads & downloads into Amazon S3



- **AWS Global Accelerator**

- Improve global application availability and performance using the AWS global network

Global Applications in AWS - Summary



- **AWS Outposts**

- Deploy Outposts Racks in your own Data Centers to extend AWS services



- **AWS WaveLength**

- Brings AWS services to the edge of the 5G networks
- Ultra-low latency applications



- **AWS Local Zones**

- Bring AWS resources (compute, database, storage, ...) closer to your users
- Good for latency-sensitive applications